

Distinction Project – Final Report

Smart Home Lighting System

Aditya Suhag (224001686) | IoT Systems – Distinction Task 3 of 3 | 4 October 2026

1. What I Ended Up With

I built the smart lighting system from my plan and ran it on real AWS. Each room has a light sensor (LDR), a motion sensor (PIR) and a smart bulb, all simulated in Node.js. Each room connects to AWS IoT Core over MQTT with TLS using its own certificate. Node-RED checks and cleans every reading. An IoT Rule triggers two Lambda functions: one logs the reading and the other decides whether the light should be on. The decision goes back down to the room and the bulb actually switches. Everything ends up in one DynamoDB table, and there's a small dashboard on top.

To show it scales, I ramped an extra 25 → 50 → 100 rooms on top of the 5 "real" ones, each sending a reading every second. I measured it from both ends: how long rooms waited for their on/off commands, and what CloudWatch says AWS was doing. Quick scorecard against the requirements in my plan:

Requirement (from my plan)	Met?	What actually happened
Each room sends light + motion every few seconds	Yes	Every 3 s from 5 rooms (load rooms every 1 s)
Works out on/off, and that changes the light	Yes	Control Lambda → MQTT command → bulb switches. Median ~400 ms end to end at normal load
Readings + decisions saved	Yes	DynamoDB history rows + a current-state row per room
Alert if a room stops sending	Yes	Node-RED watchdog alerted 12 s after the bathroom sensor "died"
Scalable: 5 → ~100 rooms, no design change	Partly	The same design ran 105 rooms with zero code changes and nothing crashed. But the account's 10-Lambda limit made it slow past ~25–50 rooms, and 6% of load readings were dropped at Node-RED (section 3)
Fast: reacts within a couple of seconds	Up to ~25 rooms	25 rooms: p50 314 ms / p95 813 ms. 100 rooms: p50 22 s
Secure: only real devices, encrypted	Yes	Cert per device, policy scoped per room (tested), least-privilege IAM, TLS
Doesn't fall over easily	Mostly	Under overload nothing crashed and Lambda dropped 0 events (just delayed them). Node-RED is a single point of failure though

2. How It Works




Figure 1: the final architecture. Solid arrows = readings going up, dashed = on/off commands coming back down.

The flow: a room publishes `{lux, motion, light, ts}` to `home/<room>/telemetry` every 3 seconds. Node-RED (on my laptop, acting as the edge/gateway) subscribes to `home/+telemetry`. It throws out anything broken (wrong types, lux out of range, room ID not matching the topic, timestamps more than 60s off, duplicates) and republishes the clean reading to `home/<room>/clean`. It also runs a per-room watchdog that raises an alert on `home/alerts/<room>` if a room goes quiet for 15 seconds. The IoT Rule `SELECT * FROM 'home/+clean'` sends each clean reading to both Lambdas at the same time.

The decision logic (control function) is still just rules, no ML:

- **Smoothing:** average this reading with the previous 2 for that room.
- **Motion + off delay:** the time of the last motion is stored, so the light stays on until there's been 30s with no motion.
- **Two thresholds (this wasn't in my plan):** the light turns on below 250 lux. But once it's on, it only turns off for brightness above 650 lux. The plan's single threshold made every light flicker, because the bulb's own light (~350 lux) made the room look "bright enough" (section 5).
- A command only gets published if the decision is different from what the room says its light is doing, so there are no pointless messages.

Data: one DynamoDB table `SmartLighting` (on-demand billing). Partition key is `roomId`, sort key is `timestamp`. Every reading is a row (raw lux, motion, smoothed lux, decision). Each room also has a special row with `timestamp = 0` holding its current state (light, lastSeen, lastMotionAt). Reading that one row is much cheaper than querying history. History rows expire after 7 days using TTL. Both functions use `UpdateItem` rather than `PutItem` because they write to the same row at the same time.

3. Scalability – The Load Test

How I tested it: I made 100 extra rooms (`load-001` to `load-100`). Each one is a proper IoT thing with its own certificate, so they go through exactly the same path as a real room. `loadtest.js` runs them across 4 worker processes and ramps `25 → 50 → 100` rooms, 2 minutes per step. Each load room sends every **1 second**, which is $3\times$ faster than my normal rooms, so 100 load rooms is about the same traffic as 300 normal ones. The 5 normal rooms kept running the whole time, plus a few minutes before and after as a baseline and cool-down. I measured two things. Each room timed how long it took from sending a reading to getting its on/off command back. I also pulled the CloudWatch metrics and graphs straight from AWS with `get-metric-widget-image`.

Load rooms	Readings/s sent	Lambda invocations / min (each fn)	Throttles / min (both fns)	Longest wait in queue	Command latency p50 / p95
0 (baseline, 5 rooms)	1.7	70–144	0	0.1 s	~0.3–1 s
25	24.4	~1,600	82 → 243	3.4 s	314 ms / 813 ms
50	49.3	~2,860	931 → 985	80 s	1.5 s / 43.5 s
100	98.7	~3,000 (flat!)	1,215 → 1,307	133 s	22.4 s / 83.8 s




Figure 2: concurrent executions. Both functions scaled out by themselves (5 → 9 each), but the whole account (red) was pinned at its limit of 10 from the moment the load test started.




Figure 3: invocations climb with each step, then flatten at ~3,000/min per function even though the rooms doubled.




Figure 4: throttles go from 0 to over 600 a minute per function. They carry on after the test while the backlog drains.




Figure 5: async event age = how long an IoT event sat in Lambda's queue before running. It climbs the whole test and peaks at ~295 s, 2 minutes after the test stopped.

What this shows:

- **Lambda scaled exactly how it's meant to, until it hit the account limit.** Once the account was at 10 concurrent executions, each function topped out at about 3,000 invocations a minute (~50 readings/s). At 100

rooms readings were arriving at ~100/s, so the rest got throttled and went into Lambda's retry queue. That's why the median latency went from 314 ms to 22 seconds.

- **Lambda didn't lose anything.** CloudWatch shows 20,615 events received vs 20,547 invoked for the log function (the rest ran just after my window), and `AsyncEventsDropped` was 0 every minute. All 6,400 throttles were retried. So overloading it made it slow, not broken.
- **But 6% of the load readings never got saved, and that was a second bottleneck I didn't expect.** 19,439 of the 20,687 load readings are in DynamoDB. At 100 rooms Node-RED's *one* MQTT connection was receiving ~100 msg/s, which is IoT Core's per-connection limit. IoT Core throttled delivery to it (`PublishOut.Throttle` = 214 at 22:05), so messages arrived late and out of order. My duplicate filter ("must be newer than the last one") then threw out 721 as "duplicate/old" and 306 as more than 60 s late, which accounts for most of the missing ones. Node-RED's CPU was fine. It was the single connection and my too-strict filter.
- **Other things held up fine.** DynamoDB on-demand took ~12,000 write units a minute without any problem. The Lambdas actually got *faster* under load (control function ~50–70 ms → ~30–40 ms average, log function ~22 → ~15 ms) because warm containers reuse their connections.
- **The load rooms hurt my real rooms (noisy neighbour).** The 5 normal rooms share the same limit, so their command latency went from under a second to an average of 12–43 s. After the test, the backlog delivered commands up to 186 s old, and their lights flapped 31 times a minute until it cleared at about 22:08.

So is it scalable? The *design* is. Going from 5 to 105 rooms needed no code or design changes: IoT Core took everything the rooms sent, Lambda scaled out automatically, and DynamoDB didn't notice. What capped it were two limits I hadn't planned for: the new account's Lambda concurrency (10) and the per-connection throughput of my single Node-RED gateway. Both are fixable with config, not a redesign (section 6), but I couldn't fully *prove* it scales past ~25–50 rooms on this account.

4. Security

Every device (and Node-RED) has its own X.509 certificate. The IoT policy uses `${iot:Connection.Thing.ThingName}`, so a device can only connect as itself, publish to its own `telemetry` topic and subscribe to its own `cmd` topic. The 100 load-test rooms got their own certs too, so the load test went through the exact same security as a real device. Each Lambda has its own IAM role with only what it needs. The log function can only `UpdateItem` on the one table. The control function can only `Query/Get/Update` that table and publish to `home/*/cmd`. Data is encrypted in transit (TLS to IoT Core on port 8883) and at rest (DynamoDB default encryption). I tested the policy with the lounge's certificate. Publishing to its own topic worked. Trying to publish to the kitchen's topic, connecting with the kitchen's client ID, and connecting with no certificate all got the connection killed by AWS (`security_test.log`).

5. What Went Wrong (and What I Learned)

- **Lights flickering (the plan's rule was wrong).** On the first full run every light switched on, saw its own glow (~350 lux), decided the room was bright, switched off, and repeated: 61 switches in about 90 seconds. Adding the second threshold (on < 250, off > 650 lux) fixed it, and the next run had 10 switches in about 4 minutes. Lesson: the sensor and the thing you're controlling affect each other, which I completely missed in the plan.
- **Rebooted devices got blocked.** My first duplicate filter in Node-RED used a sequence number. That resets to 1 when a device restarts, so everything from a restarted room got rejected as a "duplicate". I switched to the device timestamp.

- **The 10-concurrency limit** (section 3). I did spot it before the load test and flagged it in my status update. I still underestimated how fast it would bite, because I'd only thought about the *average* load and not the bursts.
- **Synchronised devices = bursts.** All 5 rooms started at the same moment and send every 3 seconds, so their readings arrive together. Even with just 5 rooms, CloudWatch shows peaks of 10 concurrent executions (5 rooms × 2 functions at once). Real devices should add a bit of random jitter to when they send.
- **Old commands after a backlog.** When the throttled events finally got processed (up to ~3 min late), the rooms acted on commands that were already out of date. Each room should ignore a command whose `basedOn` timestamp is more than a few seconds old, and the control function should skip readings that are already stale. I ran out of time to add and re-test that, so it's the first thing on my fix list.

6. Limitations and What I'd Do Next

- **Lift the concurrency limit** with a Service Quotas request (free) and re-run the same test. Everything in the design points to it scaling, but I can only *prove* it up to the limit.
- **Put SQS between the IoT Rule and the Lambdas** so each invocation handles a batch of readings. Ten readings per invocation means about a tenth of the concurrency, and the queue soaks up bursts in a controlled way rather than relying on throttle-and-retry.
- **Node-RED is a single point of failure** because it runs on my laptop. If it's down, nothing reaches the Lambdas. When I restarted it once, the lounge's history in DynamoDB has a 12-second hole (21:50:34 → 21:50:46), because those readings were simply never processed. The validation could move into the IoT Rule SQL (`WHERE lux >= 0 AND ...`), or Node-RED could run as 2+ instances using MQTT shared subscriptions. It also turned out to be a throughput limit at 100 rooms (section 3). Shared subscriptions (`$share/nodered/home/+/telemetry`) would spread the readings over several connections, and the duplicate filter should remember a short window of recent timestamps instead of only the newest one.
- **Dashboard security:** right now it uses my AWS credentials on my laptop. Properly, it would go through API Gateway with Cognito logins.
- **Multiple houses:** the topics would become `home/<houseId>/<room>/...` and the partition key `houseId#roomId`. The policy variable trick still works, so nothing else in the design has to change.
- **Real hardware:** an ESP32 with an LDR and PIR could replace `room.js` without touching anything above the MQTT connection, since it would use the same cert, topics and payload.

7. Conclusion

All the functional requirements from my plan work on real AWS: readings every few seconds, automatic on/off, everything saved, offline alerts, and per-device security that I actually tested. The scalability test was the most useful part, even though the result wasn't the perfect graph I was hoping for. The serverless design scaled as far as my account allowed. IoT Core took 100 msg/s from the rooms without an error, but Lambda hit the new-account limit of 10 at roughly 25–50 rooms, and my single Node-RED connection hit IoT Core's per-connection limit at 100 rooms. Past that point nothing crashed and Lambda didn't drop anything, but everything got very slow and 6% of readings were lost at the gateway. That's exactly the kind of thing you only find by load testing. If I did this again I'd request the quota increase on day one, add send jitter and stale-command checks to the devices, put a queue in front of the functions, and spread Node-RED over shared subscriptions.

Appendix A – Evidence Files

All of these are in the `evidence/` folder of the repo. Logs are the real terminal output from the runs (I just saved them with `tee`).

File	What it shows
<code>setup.log</code>	<code>infra/setup.sh</code> creating everything on a brand new AWS account
<code>bug_flicker_run1.log</code>	first full run with the plan's single-threshold rule – 61 light switches in ~90s (the flicker bug)
<code>house_run2.log</code>	run after the hysteresis fix, with the bathroom sensor "dying" after 60s
<code>nodered_run1.log</code> , <code>nodered_run2.log</code>	Node-RED output: offline/back-online alerts, rejected readings (incl. the duplicate-seq bug)
<code>security_test.log</code>	lounge's cert trying to publish as / connect as the kitchen, and connecting with no cert
<code>dynamodb_sample_bedroom1.json</code> , <code>dynamodb_state_row_bedroom1.json</code>	real rows from the table (history rows + the timestamp-0 state row)
<code>lambda_control_logs_sample.txt</code>	CloudWatch REPORT lines for the control function (duration, memory)
<code>loadtest_*.json</code> , <code>loadtest_console.log</code>	load test results per step (messages, commands, latency percentiles)
<code>house_during_loadtest.log</code>	the 5 normal rooms, which kept running the whole time the load test was on
<code>cloudwatch/*.png</code>	graphs pulled straight from CloudWatch with <code>get-metric-widget-image</code>
<code>cloudwatch/per_minute.json</code> , <code>metric_data.json</code>	the raw per-minute CloudWatch numbers behind the graphs
<code>nodered_flow.png</code> , <code>dashboard.png</code>	screenshots of the Node-RED editor and the dashboard

Appendix B – Code

Everything is in my local git repo `sit314_distinction` (GitHub link: *to be added*). The Node-RED `flows.json` is generated by `buildflow.js`, so I've put that script here instead of the 300-line JSON.

`simulator/room.js`

```
// One simulated room = LDR light sensor + PIR motion sensor + a smart bulb, pretending to be an ESP32.
// Connects to AWS IoT Core over MQTT/TLS with its own certificate.
// node simulator/room.js lounge -> run one room
// node simulator/room.js bathroom --die-after=60 -> sensor "dies" after 60s (to test the offline alert)
const fs = require('fs');
const mqtt = require('mqtt');
const cfg = require('../config');

const LAMP_LUX = 350; // how much the bulb adds to what the LDR sees when it's on

function daylight(hour) {
  // rough sun curve: 0 before 6am / after 7pm, peaks ~800 lux around lunchtime
  if (hour < 6 || hour > 19) return 0;
  return Math.sin(Math.PI * (hour - 6) / 13) * 800;
}

class Room {
  constructor(id, opts = {}) {
    this.id = id;
    this.windowFactor = opts.windowFactor ?? 0.5;
    this.everyMs = opts.everyMs || cfg.sendEveryMs;
    this.quiet = opts.quiet || false;
    this.dieAfterMs = opts.dieAfterMs || 0;
    this.light = 'off';
    this.occupied = Math.random() < 0.5;
    this.seq = 0;
    this.stats = { sent: 0, cmds: 0, latencies: [] };
  }

  readLux() {
    const hour = process.env.SIM_HOUR ? Number(process.env.SIM_HOUR) : new Date().getHours() + new
Date().getMinutes() / 60;
    let lux = 20 + daylight(hour) * this.windowFactor + (Math.random() - 0.5) * 60; // 20 lux of background + noise
    if (this.light === 'on') lux += LAMP_LUX;
    return Math.max(0, Math.round(lux));
  }

  readMotion() {
    // people come and go randomly; while someone's in the room the PIR doesn't fire every time (sitting still etc.)
    if (this.occupied && Math.random() < 0.05) this.occupied = false;
    else if (!this.occupied && Math.random() < 0.04) this.occupied = true;
    return this.occupied && Math.random() < 0.6;
  }

  start() {
    const c = cfg.certDir;
    this.client = mqtt.connect(`mqtt://${cfg.endpoint}:8883`, {
      clientId: this.id,
      key: fs.readFileSync(`${c}/${this.id}.private.key`),
      cert: fs.readFileSync(`${c}/${this.id}.cert.pem`),
      ca: fs.readFileSync(`${c}/AmazonRootCA1.pem`),
      reconnectPeriod: 2000,
    });
    this.client.on('connect', () => {
      this.log('connected to IoT Core');
      this.client.subscribe(`home/${this.id}/cmd`, { qos: 1 });
    });
    this.client.on('error', e => this.log('mqtt error: ' + e.message, true));
    this.client.on('message', (_t, buf) => this.onCommand(JSON.parse(buf)));

    this.timer = setInterval(() => this.sendReading(), this.everyMs);
    if (this.dieAfterMs) setTimeout(() => { clearInterval(this.timer); this.log('*** sensor died (stopped sending
***', true); }, this.dieAfterMs);
  }

  sendReading() {
    if (!this.client.connected) return;
    const msg = { roomId: this.id, ts: Date.now(), seq: ++this.seq, lux: this.readLux(), motion: this.readMotion(),
light: this.light };
    this.client.publish(`home/${this.id}/telemetry`, JSON.stringify(msg), { qos: 1 });
    this.stats.sent++;
  }
}
```



```

    this.log(`lux=${String(msg.lux).padStart(4)} motion=${msg.motion ? 'Y' : '-'} light=${this.light}`);
  }

  onCommand(cmd) {
    this.stats.cmds++;
    const latency = cmd.basedOn ? Date.now() - cmd.basedOn : null; // time from the reading being sent to the
    command arriving
    if (latency != null) this.stats.latencies.push(latency);
    if (cmd.state !== this.light) {
      this.light = cmd.state;
      this.log(`>>> light switched ${cmd.state.toUpperCase()} (${cmd.reason}) latency=${latency}ms`, true);
    }
  }

  stop() { clearInterval(this.timer); this.client?.end(true); }

  log(s, important = false) {
    if (this.quiet && !important) return;
    console.log(`[${new Date().toLocaleTimeString('en-AU', { hour12: false })}] ${this.id.padEnd(9)} ${s}`);
  }
}

module.exports = { Room };

if (require.main === module) {
  const id = process.argv[2];
  const room = cfg.rooms.find(r => r.id === id) || { id, windowFactor: 0.5 };
  const die = process.argv.find(a => a.startsWith('--die-after='));
  new Room(id, { windowFactor: room.windowFactor, dieAfterMs: die ? Number(die.split('=')[1]) * 1000 : 0 }).start();
}

```

simulator/house.js

```

// starts all 5 rooms, one Node.js process each (like 5 separate ESP32s)
// extra args get passed through, e.g. node simulator/house.js --die=bathroom:60
const { fork } = require('child_process');
const cfg = require('../config');

const die = (process.argv.find(a => a.startsWith('--die=')) || '').replace('--die=', '').split(':');

for (const r of cfg.rooms) {
  const args = [r.id];
  if (die[0] === r.id) args.push(`--die-after=${die[1]}`);
  fork(__dirname + '/room.js', args);
}

console.log(`started ${cfg.rooms.length} rooms: ${cfg.rooms.map(r => r.id).join(', ')}`);

```

lambdas/log-reading/index.mjs

```

// Lambda 1: Log Function - saves every cleaned reading into DynamoDB + bumps the room's lastSeen.
// Triggered by the IoT Rule on topic home/+/clean (Node-RED publishes there after checking the data).
import { DynamoDBClient } from '@aws-sdk/client-dynamodb';
import { DynamoDBDocumentClient, UpdateCommand } from '@aws-sdk/lib-dynamodb';

const db = DynamoDBDocumentClient.from(new DynamoDBClient({}));
const TABLE = process.env.TABLE_NAME;
const KEEP_DAYS = 7; // TTL so the history doesn't just grow forever

export const handler = async (r) => {
  const expiresAt = Math.floor(r.ts / 1000) + KEEP_DAYS * 86400;

  // UpdateItem instead of PutItem because the control function writes to the same row at the same time,
  // and Put would wipe out whatever it already wrote
  await db.send(new UpdateCommand({
    TableName: TABLE,
    Key: { roomId: r.roomId, timestamp: r.ts },
    UpdateExpression: 'SET rawLux = :lux, motionDetected = :m, reportedLight = :l, expiresAt = :e',
    ExpressionAttributeValues: { ':lux': r.lux, ':m': r.motion, ':l': r.light, ':e': expiresAt },
  }));

  // timestamp 0 = the "current state" row for each room (one table, so this is my little shortcut)
  await db.send(new UpdateCommand({
    TableName: TABLE,
    Key: { roomId: r.roomId, timestamp: 0 },
    UpdateExpression: 'SET lastSeen = :ts, lastRawLux = :lux',
    ExpressionAttributeValues: { ':ts': r.ts, ':lux': r.lux },
  }));

  return { saved: `${r.roomId}#${r.ts}` };
};

```


lambdas/control-light/index.mjs

```
// Lambda 2: Control Function - decides if the light should be on or off and sends the command back down.
// Triggered by the same IoT Rule as the log function (home+/clean).
import { DynamoDBClient } from '@aws-sdk/client-dynamodb';
import { DynamoDBDocumentClient, QueryCommand, GetCommand, UpdateCommand } from '@aws-sdk/lib-dynamodb';
import { IoTDataPlaneClient, PublishCommand } from '@aws-sdk/client-iot-data-plane';

const db = DynamoDBDocumentClient.from(new DynamoDBClient({}));
const iot = new IoTDataPlaneClient({ endpoint: `https://${process.env.IOT_ENDPOINT}` });
const TABLE = process.env.TABLE_NAME;

const DARK = 250;           // lux - below this the room counts as dark
const BRIGHT = 650;         // lux - once the light is ON it has to be brighter than this before it turns off for
                             // brightness.
                             // the bulb itself adds ~350 lux to what the sensor sees, so without this gap the light
                             // turned itself off straight away and flickered every few seconds (found this on my
                             // first test run)
const OFF_DELAY_MS = 30000; // wait 30s after the last motion before switching off

export const handler = async (r) => {
  // smoothing: average this reading with the previous 2.
  // (the log function runs in parallel so this reading might not be in the table yet - that's why I add r.lux
  myself)
  const prev = await db.send(new QueryCommand({
    TableName: TABLE,
    KeyConditionExpression: 'roomId = :id AND #t BETWEEN :from AND :to',
    ExpressionAttributeNames: { '#t': 'timestamp' },
    ExpressionAttributeValues: { ':id': r.roomId, ':from': 1, ':to': r.ts - 1 },
    ScanIndexForward: false, // newest first
    Limit: 2,
  }));
  const luxes = [r.lux, ...prev.Items.map(i => i.rawLux).filter(x => x != null)];
  const avgLux = Math.round(luxes.reduce((a, b) => a + b, 0) / luxes.length);

  const state = (await db.send(new GetCommand({ TableName: TABLE, Key: { roomId: r.roomId, timestamp: 0 } }))).Item
  || {};
  const lastMotionAt = r.motion ? r.ts : (state.lastMotionAt || 0);
  const motionRecent = r.ts - lastMotionAt < OFF_DELAY_MS;

  // the actual rule: someone's been here recently AND it's dark -> on, otherwise off
  // (with two thresholds so the light doesn't count its own glow as daylight)
  const threshold = r.light === 'on' ? BRIGHT : DARK;
  const isDark = avgLux < threshold;
  const decision = motionRecent && isDark ? 'on' : 'off';
  const reason = !motionRecent ? 'no motion for 30s' : isDark ? `dark (avg ${avgLux} < ${threshold} lux) + motion` :
  `bright enough (avg ${avgLux} > ${threshold} lux)`;

  // only send a command if the light isn't already in the right state
  if (decision !== r.light) {
    await iot.send(new PublishCommand({
      topic: `home/${r.roomId}/cmd`,
      qos: 1,
      payload: JSON.stringify({ state: decision, reason, avgLux, basedOn: r.ts }),
    }));
  }

  await db.send(new UpdateCommand({
    TableName: TABLE,
    Key: { roomId: r.roomId, timestamp: r.ts },
    UpdateExpression: 'SET lightLevel = :avg, lightState = :s',
    ExpressionAttributeValues: { ':avg': avgLux, ':s': decision },
  }));
  await db.send(new UpdateCommand({
    TableName: TABLE,
    Key: { roomId: r.roomId, timestamp: 0 },
    UpdateExpression: 'SET lightState = :s, lightLevel = :avg, lastMotionAt = :lm, motionDetected = :m, reason = :why',
    ExpressionAttributeValues: { ':s': decision, ':avg': avgLux, ':lm': lastMotionAt, ':m': r.motion, ':why': reason },
  }));

  return { roomId: r.roomId, decision, avgLux, sentCommand: decision !== r.light };
};
```

nodered/buildflow.js (generates the Node-RED flows.json)

```
// generates flows.json for Node-RED (easier than dragging 15 nodes around by hand every time I change something)
// run: node nodered/buildflow.js then node-red --userDir nodered
const fs = require('fs');
const path = require('path');
const cfg = require('../config');
```

```

const certs = path.resolve(__dirname, '../certs');

const tab = 'tab_lighting';
const n = (id, type, x, y, props = {}, wires = []) => ({ id, type, z: tab, x, y, wires, ...props });

const validateFn = `
// checks each reading before it goes anywhere near the cloud functions
const room = msg.topic.split('/')[1];
const r = msg.payload;
const problems = [];

if (typeof r !== 'object' || r === null) problems.push('not valid JSON');
else {
  if (r.roomId !== room) problems.push('roomId does not match topic');
  if (typeof r.lux !== 'number' || isNaN(r.lux) || r.lux < 0 || r.lux > 100000) problems.push('lux missing/out of range');
  if (typeof r.motion !== 'boolean') problems.push('motion should be true/false');
  if (r.light !== 'on' && r.light !== 'off') problems.push('light state should be on/off');
  if (typeof r.ts !== 'number' || Math.abs(Date.now() - r.ts) > 60000) problems.push('timestamp missing or more than 60s off');
}

// QoS 1 can deliver the same message twice, so drop anything that isn't newer than the last one from that room.
// (I used the seq number at first, but it resets to 1 when a device restarts, so after a reboot EVERYTHING
// got thrown away as a duplicate. The device timestamp doesn't have that problem.)
const seen = context.get('lastTs') || {};
if (!problems.length && seen[room] && r.ts <= seen[room]) problems.push('duplicate/old (ts ' + r.ts + ')');

const stats = context.get('stats') || { ok: 0, bad: 0 };
if (problems.length) {
  stats.bad++; context.set('stats', stats);
  node.status({ fill: 'yellow', shape: 'dot', text: stats.ok + ' ok / ' + stats.bad + ' rejected' });
  return [null, { topic: room, payload: { room, problems, original: r } }];
}

seen[room] = r.ts; context.set('lastTs', seen);
stats.ok++; context.set('stats', stats);
if (stats.ok % 10 === 0) node.status({ fill: 'green', shape: 'dot', text: stats.ok + ' ok / ' + stats.bad + ' rejected' });

msg.topic = 'home/' + room + '/clean';
msg.payload = { roomId: room, ts: r.ts, lux: Math.round(r.lux), motion: r.motion, light: r.light, receivedAt: Date.now() };
return [msg, null];`;

const onlineFn = `
// if a room that was flagged offline starts talking again, say so
const room = msg.topic.split('/')[1];
const offline = flow.get('offline') || {};
if (offline[room]) {
  delete offline[room]; flow.set('offline', offline);
  node.warn('room ' + room + ' is back online');
  node.send([null, { topic: 'home/alerts/' + room, payload: { roomId: room, alert: 'back-online', at: Date.now() } }]);
}
return [msg, null];`;

const alertFn = `
// the trigger node only lets a message through here if a room has gone quiet for ${cfg.offlineAfterMs / 1000}s
const room = msg.topic.split('/')[1];
const offline = flow.get('offline') || {};
offline[room] = Date.now(); flow.set('offline', offline);
const lastSeen = msg.payload && msg.payload.ts;
msg.topic = 'home/alerts/' + room;
msg.payload = { roomId: room, alert: 'offline', lastSeen, at: Date.now(), message: 'no data from ' + room + ' for ${cfg.offlineAfterMs / 1000}s - sensor probably dead' };
node.status({ fill: 'red', shape: 'ring', text: room + ' OFFLINE' });
return msg;`;

const flows = [
  { id: tab, type: 'tab', label: 'Smart Lighting', info: 'room sensors -> check & clean -> AWS (IoT rule -> Lambdas)' },
  { id: 'broker', type: 'mqtt-broker', name: 'AWS IoT Core', broker: cfg.endpoint, port: '8883', clientid: 'nodered-gateway',
    usetls: true, tls: 'tls', protocolVersion: '4', keepalive: '60', cleansession: true, autoConnect: true,
    birthTopic: '', closeTopic: '', willTopic: '' },
  { id: 'tls', type: 'tls-config', name: 'nodered-gateway cert', cert: `${certs}/nodered-gateway.cert.pem`,
    key: `${certs}/nodered-gateway.private.key`, ca: `${certs}/AmazonRootCA1.pem`, certname: '', keyname: '',
    caname: '',
    servername: '', verifyservercert: true },
  { id: 'c1', type: 'comment', z: tab, name: '1) readings from every room -> check & clean -> back to IoT Core on home/<room>/clean (IoT Rule sends them to the Lambdas)', x: 470, y: 40, wires: [] },
  n('in', 'mqtt in', 140, 100, { name: 'home+/telemetry', topic: 'home+/telemetry', qos: '1', datatype: 'json',

```

```

broker: 'broker', nl: false, rap: true, rh: 0, inputs: 0 }, [['online']]),
  n('online', 'function', 360, 100, { name: 'back online?', func: onlineFn, outputs: 2 }, [['validate', 'watchdog'],
['alertOut', 'alertDbg']]),
  n('validate', 'function', 580, 100, { name: 'check & clean', func: validateFn, outputs: 2 }, [['cleanOut'],
['rejDbg']]),
  n('cleanOut', 'mqtt out', 830, 80, { name: 'home/<room>/clean', topic: '', qos: '1', retain: 'false', broker:
'broker' }),
  n('rejDbg', 'debug', 830, 140, { name: 'REJECTED reading', active: true, tosidebar: true, console: true, complete:
'payload', targetType: 'msg' })),

  { id: 'c2', type: 'comment', z: tab, name: '2) watchdog - if a room goes quiet for 15s, raise an offline alert',
x: 330, y: 220, wires: [] },
  n('watchdog', 'trigger', 380, 280, { name: 'wait 15s (reset on each msg, per room)', op1: '', op2: '', op1type:
'nul', op2type: 'payl',
  duration: String(cfg.offlineAfterMs / 1000), extend: true, overrideDelay: false, units: 's', reset: '', bytopic:
'topic', topic: 'topic', outputs: 1 }, [['alert']]),
  n('alert', 'function', 640, 280, { name: 'build offline alert', func: alertFn, outputs: 1 }, [['alertOut',
'alertDbg']]),
  n('alertOut', 'mqtt out', 870, 260, { name: 'home/alerts/<room>', topic: '', qos: '1', retain: 'false', broker:
'broker' }),
  n('alertDbg', 'debug', 860, 320, { name: 'ALERT', active: true, tosidebar: true, console: true, complete:
'payload', targetType: 'msg' })),

  { id: 'c3', type: 'comment', z: tab, name: '3) test: push some broken readings through the checker', x: 290, y:
380, wires: [] },
  n('bad1', 'inject', 160, 430, { name: 'bad lux + motion', props: [{ p: 'payload' }, { p: 'topic', vt: 'str' }],
topic: 'home/lounge/telemetry',
  payload: '{"roomId":"lounge","ts":0,"lux":-50,"motion":"yes","light":"on"}', payloadType: 'json', repeat: '',
once: false }, [['validate']]),
  n('bad2', 'inject', 170, 480, { name: 'wrong room in topic', props: [{ p: 'payload' }, { p: 'topic', vt: 'str' }],
topic: 'home/kitchen/telemetry',
  payload: '{"roomId":"lounge","ts":1,"lux":100,"motion":true,"light":"off"}', payloadType: 'json', repeat: '',
once: false }, [['validate']]),
  ];

fs.writeFileSync(__dirname + '/flows.json', JSON.stringify(flows, null, 2));
console.log('wrote nodered/flows.json with', flows.length, 'nodes');

```

loadtest/loadtest.js

```

// Load test: ramps up simulated rooms in steps (e.g. 25 -> 50 -> 100) so I can see Lambda scale up, then stops so I
can see it scale down.
// Every load room is a proper device with its own cert, it's just that several rooms share one Node.js process
// (100 separate node processes = ~4GB of RAM on my laptop, so I split them over a few worker processes instead).
//
// node loadtest/loadtest.js --steps=25,50,100 --step-secs=120 --every=1000 --workers=4
const { fork } = require('child_process');
const fs = require('fs');

const arg = (k, d) => (process.argv.find(a => a.startsWith(`--${k}=`)) || `=${d}`).split('=')[1];

if (process.argv[2] === '--worker') return worker();

const steps = arg('steps', '25,50,100').split(',').map(Number);
const stepSecs = Number(arg('step-secs', 120));
const every = Number(arg('every', 1000));
const workers = Number(arg('workers', 4));
const total = Math.max(...steps);
const startedAt = Date.now();

console.log(`load test: steps ${steps.join(' -> ')} rooms, ${stepSecs}s each, every room sends every ${every}ms,
${workers} worker processes`);
console.log(`started ${new Date(startedAt).toISOString()}`);

// work out when each room should start (room i joins at the first step that includes it)
const rooms = [];
for (let i = 0; i < total; i++) {
  const step = steps.findIndex(s => i < s);
  rooms.push({ id: `load-${String(i + 1).padStart(3, '0')}`, startAfterMs: step * stepSecs * 1000 + (i % 25) * 200,
step });
}

const results = [];
let done = 0;
for (let w = 0; w < workers; w++) {
  const mine = rooms.filter((_, i) => i % workers === w);
  const child = fork(__filename, ['--worker']);
  child.send({ rooms: mine, every, nSteps: steps.length, stopAtMs: steps.length * stepSecs * 1000 });
  child.on('message', m => {
    if (m.type === 'progress') return;
    results.push(...m.rooms);
    if (++done === workers) report();
  });
};

```

```

}

// print a line every 10s so I can see it's going
const tick = setInterval(() => {
  const secs = Math.round((Date.now() - startedAt) / 1000);
  const step = Math.min(Math.floor(secs / stepSecs), steps.length - 1);
  console.log(` t=${secs}s step ${step + 1}/${steps.length} (${steps[step]} rooms)`);
}, 10000);

function pct(arr, p) { if (!arr.length) return null; const s = [...arr].sort((a, b) => a - b); return
s[Math.min(s.length - 1, Math.floor(p / 100 * s.length))]; }

function report() {
  clearInterval(tick);
  const endedAt = Date.now();
  const summary = { startedAt: new Date(startedAt).toISOString(), endedAt: new Date(endedAt).toISOString(), steps,
stepSecs, everyMs: every, workers, perStep: [] };
  steps.forEach((s, k) => {
    const lat = results.flatMap(r => r.latencies.filter(l => l.step === k).map(l => l.ms));
    const sent = results.reduce((a, r) => a + r.sentPerStep[k], 0);
    summary.perStep.push({ rooms: s, messagesSent: sent, msgPerSec: +(sent / stepSecs).toFixed(1), commandsReceived:
lat.length,
    latencyMs: { p50: pct(lat, 50), p95: pct(lat, 95), p99: pct(lat, 99), max: lat.length ? Math.max(...lat) :
null } });
  });
  summary.totalSent = results.reduce((a, r) => a + r.sent, 0);
  summary.totalCommands = results.reduce((a, r) => a + r.cmds, 0);
  summary.connectErrors = results.reduce((a, r) => a + r.errors, 0);

  console.log('\n=== results ===');
  console.table(summary.perStep.map(p => ({ rooms: p.rooms, 'msgs sent': p.messagesSent, 'msg/s': p.msgPerSec, 'cmds
back': p.commandsReceived,
  'p50 ms': p.latencyMs.p50, 'p95 ms': p.latencyMs.p95, 'p99 ms': p.latencyMs.p99, 'max ms': p.latencyMs.max })));
  console.log(`total sent ${summary.totalSent}, commands received ${summary.totalCommands}, mqtt errors
${summary.connectErrors}`);
  const out = `_${dirname}/../evidence/loadtest_${new Date(startedAt).toISOString().slice(0, 16).replace(/:/g,
'')}.json`;
  fs.writeFileSync(out, JSON.stringify(summary, null, 2));
  console.log('saved', out);
  process.exit(0);
}

function worker() {
  const { Room } = require('../simulator/room');
  process.on('message', ({ rooms, every, nSteps, stopAtMs }) => {
    const t0 = Date.now();
    const live = rooms.map(r => {
      const room = new Room(r.id, { windowFactor: Math.random(), everyMs: every, quiet: true });
      room.sentPerStep = {}; room.errors = 0;
      // wrap so I know which step each message/latency belongs to
      const send = room.sendReading.bind(room);
      room.sendReading = () => { const before = room.stats.sent; send(); if (room.stats.sent > before) { const k =
step(); room.sentPerStep[k] = (room.sentPerStep[k] || 0) + 1; } };
      const onCmd = room.onCommand.bind(room);
      room.onCommand = cmd => { onCmd(cmd); if (cmd.basedOn) room.lat.push({ step: step(), ms: Date.now() -
cmd.basedOn }); };
      room.lat = [];
      setTimeout(() => { room.start(); room.client.on('error', () => room.errors++); }, r.startAfterMs);
      return room;
    });
    function step() { return Math.min(Math.floor((Date.now() - t0) / (stopAtMs / nSteps)), nSteps - 1); }

    setTimeout(() => {
      live.forEach(r => r.stop());
      process.send({ type: 'done', rooms: live.map(r => ({ id: r.id, sent: r.stats.sent, cmds: r.stats.cmds, errors:
r.errors, latencies: r.lat,
      sentPerStep: Array.from({ length: nSteps }, (_, k) => r.sentPerStep[k] || 0) })));
      setTimeout(() => process.exit(0), 500);
    }, stopAtMs);
  });
}

```

loadtest/cloudwatch-graphs.js

```

// pulls the CloudWatch graphs for a load test run straight from AWS (GetMetricWidgetImage) + the raw numbers
(GetMetricData)
// node loadtest/cloudwatch-graphs.js evidence/loadtest_XXXX.json
const { execFileSync } = require('child_process');
const fs = require('fs');

const run = JSON.parse(fs.readFileSync(process.argv[2], 'utf8'));
const t0 = new Date(run.startedAt), t1 = new Date(run.endedAt);
const start = new Date(t0 - 4 * 60000).toISOString(); // 4 min of baseline before

```

```

const end = new Date(+t1 + 5 * 60000).toISOString(); // 5 min of cool-down after
const out = __dirname + '/../evidence/cloudwatch';
fs.mkdirSync(out, { recursive: true });

// vertical lines where each step started/ended
const marks = run.steps.map((s, k) => ({ label: `${s} rooms`, value: new Date(+t0 + k * run.stepSecs * 1000).toISOString(), color: '#7b1020' }));
marks.push({ label: 'load test stops', value: run.endedAt, color: '#555555' });

const fns = ['lighting-log-reading', 'lighting-control-light'];
const lambda = (metric) => fns.map(f => ['AWS/Lambda', metric, 'FunctionName', f, { label: f.replace('lighting-', '') }]);

const graphs = {
  lambda_concurrency: { title: 'Lambda concurrent executions (max per minute)', stat: 'Maximum', metrics: [...lambda('ConcurrentExecutions'), ['AWS/Lambda', 'ConcurrentExecutions', { label: 'whole account (limit = 10)', color: '#d62728' }]] },
  lambda_invocations: { title: 'Lambda invocations per minute', stat: 'Sum', metrics: lambda('Invocations') },
  lambda_duration: { title: 'Lambda duration (ms, average)', stat: 'Average', metrics: lambda('Duration') },
  lambda_throttles: { title: 'Lambda throttles per minute', stat: 'Sum', metrics: [...lambda('Throttles'), ['AWS/Lambda', 'Throttles', { label: 'whole account' }]] },
  lambda_async_age: { title: 'Async event age (ms) - how long IoT events waited before Lambda ran them', stat: 'Maximum', metrics: lambda('AsyncEventAge') },
  iot_messages: { title: 'IoT Core messages per minute', stat: 'Sum', metrics: [
    ['AWS/IoT', 'PublishIn.Success', 'Protocol', 'MQTT', { label: 'published in (rooms + Node-RED)' }],
    ['AWS/IoT', 'PublishOut.Success', 'Protocol', 'MQTT', { label: 'delivered out (to Node-RED + rooms)' }],
    ['AWS/IoT', 'TopicMatch', 'RuleName', 'LightingReadings', { label: 'IoT Rule matches' }]] },
  dynamodb_writes: { title: 'DynamoDB consumed write capacity per minute', stat: 'Sum', metrics: [['AWS/DynamoDB', 'ConsumedWriteCapacityUnits', 'TableName', 'SmartLighting', { label: 'write units' }],
    ['AWS/DynamoDB', 'ConsumedReadCapacityUnits', 'TableName', 'SmartLighting', { label: 'read units' }]] },
};

for (const [name, g] of Object.entries(graphs)) {
  const widget = { width: 900, height: 330, start, end, period: 60, stat: g.stat, view: 'timeSeries', title: g.title, metrics: g.metrics,
    annotations: { vertical: marks }, yAxis: { left: { min: 0 } }, legend: { position: 'bottom' }, timezone: '+1100' };
  const b64 = execFileSync('aws', ['cloudwatch', 'get-metric-widget-image', '--metric-widget', JSON.stringify(widget), '--output', 'text', '--query', 'MetricWidgetImage']).toString();
  fs.writeFileSync(`${out}/${name}.png`, Buffer.from(b64, 'base64'));
  console.log('saved', name + '.png');
}

// raw numbers (per minute) so I can put actual values in the report, not just squint at graphs
const queries = [];
const q = (id, ns, metric, dims, stat) => queries.push({ Id: id, MetricStat: { Metric: { Namespace: ns, MetricName: metric, Dimensions: dims }, Period: 60, Stat: stat } });
fns.forEach((f, i) => {
  const d = [{ Name: 'FunctionName', Value: f }];
  q(`conc${i}`, 'AWS/Lambda', 'ConcurrentExecutions', d, 'Maximum');
  q(`inv${i}`, 'AWS/Lambda', 'Invocations', d, 'Sum');
  q(`dur${i}`, 'AWS/Lambda', 'Duration', d, 'Average');
  q(`p99${i}`, 'AWS/Lambda', 'Duration', d, 'p99');
  q(`thr${i}`, 'AWS/Lambda', 'Throttles', d, 'Sum');
  q(`age${i}`, 'AWS/Lambda', 'AsyncEventAge', d, 'Maximum');
  q(`drop${i}`, 'AWS/Lambda', 'AsyncEventsDropped', d, 'Sum');
  q(`recv${i}`, 'AWS/Lambda', 'AsyncEventsReceived', d, 'Sum');
});
q('accConc', 'AWS/Lambda', 'ConcurrentExecutions', [], 'Maximum');
q('pubIn', 'AWS/IoT', 'PublishIn.Success', [{ Name: 'Protocol', Value: 'MQTT' }], 'Sum');
q('ruleMatch', 'AWS/IoT', 'TopicMatch', [{ Name: 'RuleName', Value: 'LightingReadings' }], 'Sum');
q('wcu', 'AWS/DynamoDB', 'ConsumedWriteCapacityUnits', [{ Name: 'TableName', Value: 'SmartLighting' }], 'Sum');
const data = JSON.parse(execFileSync('aws', ['cloudwatch', 'get-metric-data', '--start-time', start, '--end-time', end, '--metric-data-queries', JSON.stringify(queries), '--output', 'json']).toString());
fs.writeFileSync(`${out}/metric_data.json`, JSON.stringify(data, null, 2));

// one row per minute
const rows = {};
for (const r of data.MetricDataResults) r.Timestamps.forEach((t, i) => { (rows[t] ||= {})[r.Id] = Math.round(r.Values[i] * 10) / 10; });
const table = Object.keys(rows).sort().map(t => ({ minute: new Date(t).toLocaleTimeString('en-AU', { hour12: false, timeZone: 'Australia/Melbourne' }).slice(0, 5), ...rows[t] }));
console.table(table);
fs.writeFileSync(`${out}/per_minute.json`, JSON.stringify(table, null, 2));

```

simulator/security-test.js

```

// checks the IoT policy actually stops a device pretending to be another room
// test 1: lounge's cert tries to publish to kitchen's topic -> AWS should kick it off
// test 2: lounge's cert tries to connect with clientId "kitchen" -> AWS should refuse the connection
// test 3: no certificate at all -> TLS handshake should fail
const fs = require('fs');
const mqtt = require('mqtt');

```

```

const cfg = require('../config');
const c = cfg.certDir;
const lounge = { key: fs.readFileSync(`${c}/lounge.private.key`), cert: fs.readFileSync(`${c}/lounge.cert.pem`), ca:
fs.readFileSync(`${c}/AmazonRootCA1.pem`) };
const url = `mqtt://${cfg.endpoint}:8883`;
const log = s => console.log(`[${new Date().toLocaleTimeString('en-AU', { hour12: false })}] ${s}`);

function attempt(name, opts, after) {
  return new Promise(done => {
    log(`--- ${name}`);
    const cl = mqtt.connect(url, { ...opts, reconnectPeriod: 0, connectTimeout: 5000 });
    let over = false;
    const finish = msg => { if (over) return; over = true; log(msg); cl.end(true); setTimeout(done, 300); };
    cl.on('connect', () => { log('connected'); after ? after(cl, finish) : finish('result: connected OK'); });
    cl.on('error', e => finish('result: error -> ' + e.message));
    cl.on('close', () => finish('result: connection closed by AWS'));
  });
}

(async () => {
  await attempt('control: lounge cert, clientId lounge, publish to home/lounge/telemetry', { ...lounge, clientId:
'lounge' }, (cl, finish) => {
    cl.publish('home/lounge/telemetry', JSON.stringify({ roomId: 'lounge', ts: Date.now(), lux: 10, motion: false,
light: 'off' }), { qos: 1 },
    e => finish(e ? 'result: publish failed ' + e.message : 'result: publish accepted (PUBACK received) -
expected')));
    await attempt('test 1: lounge cert publishes to home/kitchen/telemetry', { ...lounge, clientId: 'lounge' }, (cl,
finish) => {
    cl.publish('home/kitchen/telemetry', JSON.stringify({ roomId: 'kitchen', ts: Date.now(), lux: 999, motion: true,
light: 'off' }), { qos: 1 },
    e => finish(e ? 'result: publish failed ' + e.message : 'result: publish accepted ?!')));
    await attempt('test 2: lounge cert connects as clientId "kitchen"', { ...lounge, clientId: 'kitchen' });
    await attempt('test 3: no client certificate', { ca: lounge.ca, clientId: 'lounge' });
    process.exit(0);
  })();
})();

```

dashboard/server.js

```

// tiny dashboard - reads each room's current state row (timestamp 0) + recent history from DynamoDB
// runs on my laptop with my AWS creds (fine for a demo, see report for what I'd do properly)
const express = require('express');
const { DynamoDBClient } = require('@aws-sdk/client-dynamodb');
const { DynamoDBDocumentClient, BatchGetCommand, QueryCommand } = require('@aws-sdk/lib-dynamodb');
const cfg = require('../config');

const db = DynamoDBDocumentClient.from(new DynamoDBClient({ region: cfg.region }));
const app = express();
app.use(express.static(__dirname + '/public'));

app.get('/api/rooms', async (_req, res) => {
  const out = await db.send(new BatchGetCommand({ RequestItems: { [cfg.table]: { Keys: cfg.rooms.map(r => ({ roomId:
r.id, timestamp: 0 } )) } } }));
  const rows = out.Responses[cfg.table];
  res.json(cfg.rooms.map(r => {
    const s = rows.find(x => x.roomId === r.id) || {};
    return { ...s, roomId: r.id, online: !!s.lastSeen && Date.now() - s.lastSeen < cfg.offlineAfterMs };
  }));
});

app.get('/api/history/:room', async (req, res) => {
  const out = await db.send(new QueryCommand({
    TableName: cfg.table, KeyConditionExpression: 'roomId = :r AND #t > :z',
    ExpressionAttributeNames: { '#t': 'timestamp' }, ExpressionAttributeValues: { ':r': req.params.room, ':z': 0 },
    ScanIndexForward: false, Limit: 20,
  }));
  res.json(out.Items);
});

app.listen(3000, () => console.log('dashboard on http://localhost:3000'));

```

config.js

```

// shared settings for the simulator, load test and dashboard
require('dotenv').config({ path: __dirname + '/.env', quiet: true });

module.exports = {
  region: process.env.AWS_REGION || 'ap-southeast-2',
  endpoint: process.env.IOT_ENDPOINT, // xxxx-ats.iot.ap-southeast-2.amazonaws.com
  table: process.env.TABLE_NAME || 'SmartLighting',
  certDir: __dirname + '/certs',

  // the 5 rooms in "my house". windowFactor = how much daylight gets in (0 = no window)

```

```
rooms: [
  { id: 'lounge',   windowFactor: 1.0 },
  { id: 'kitchen', windowFactor: 0.8 },
  { id: 'bedroom1', windowFactor: 0.6 },
  { id: 'bedroom2', windowFactor: 0.5 },
  { id: 'bathroom', windowFactor: 0.0 },
],

sendEveryMs: 3000,    // each room reports every 3s
offlineAfterMs: 15000, // no message for 15s = probably a dead sensor
};
```


Appendix C – Configuration (AWS scripts, policies, IoT rule)

REGION / ACCOUNT get filled in by `setup.sh` when it runs.

infra/setup.sh

```
#!/bin/bash
# sets up everything on AWS for the smart lighting project (run once)
set -e
cd "$(dirname "$0")"
REGION=${AWS_REGION:-ap-southeast-2}
export AWS_REGION=$REGION AWS_PAGER=""
ACCOUNT=$(aws sts get-caller-identity --query Account --output text)
ENDPOINT=$(aws iot describe-endpoint --endpoint-type iot:Data-ATS --query endpointAddress --output text)
echo "account $ACCOUNT, region $REGION, endpoint $ENDPOINT"

# fill in REGION/ACCOUNT in the policy files
mkdir -p build
for f in policies/*.json; do sed "s/REGION/$REGION/g; s/ACCOUNT/$ACCOUNT/g" "$f" > "build/${basename "$f"}"; done

printf "AWS_REGION=%s\nIOT_ENDPOINT=%s\nTABLE_NAME=SmartLighting\n" "$REGION" "$ENDPOINT" > ../.env

echo "== 1. DynamoDB table"
aws dynamodb create-table --table-name SmartLighting \
  --attribute-definitions AttributeName=roomId,AttributeType=S AttributeName=timestamp,AttributeType=N \
  --key-schema AttributeName=roomId,KeyType=HASH AttributeName=timestamp,KeyType=RANGE \
  --billing-mode PAY_PER_REQUEST >/dev/null
aws dynamodb wait table-exists --table-name SmartLighting
aws dynamodb update-time-to-live --table-name SmartLighting \
  --time-to-live-specification "Enabled=true, AttributeName=expiresAt" >/dev/null

echo "== 2. IAM roles (one each, only what that function needs)"
for fn in log control; do
  aws iam create-role --role-name lighting-$fn-role \
    --assume-role-policy-document file://build/lambda-trust.json >/dev/null
  aws iam put-role-policy --role-name lighting-$fn-role --policy-name lighting-$fn-access \
    --policy-document file://build/$fn-role-policy.json
done
echo "  waiting for IAM to catch up..."; sleep 12

echo "== 3. Lambda functions"
deploy() { # name folder role
  (cd ../lambdas/$2 && zip -qj ../../infra/build/$2.zip index.mjs)
  aws lambda create-function --function-name $1 --runtime nodejs22.x --handler index.handler \
    --role arn:aws:iam::$ACCOUNT:role/$3 --zip-file fileb://build/$2.zip \
    --memory-size 256 --timeout 10 \
    --environment "Variables={TABLE_NAME=SmartLighting,IOT_ENDPOINT=$ENDPOINT}" >/dev/null
  aws lambda add-permission --function-name $1 --statement-id iot-rule-invoke \
    --action lambda:InvokeFunction --principal iot.amazonaws.com \
    --source-arn arn:aws:iot:$REGION:$ACCOUNT:rule/LightingReadings >/dev/null
  echo "  $1 deployed"
}
deploy lighting-log-reading log-reading lighting-log-role
deploy lighting-control-light control-light lighting-control-role

echo "== 4. IoT Core policies, things and certificates"
aws iot create-policy --policy-name SmartLightingDevicePolicy --policy-document file://build/device-policy.json
>/dev/null
aws iot create-policy --policy-name SmartLightingGatewayPolicy --policy-document file://build/gateway-policy.json
>/dev/null
curl -s https://www.amazontrust.com/repository/AmazonRootCA1.pem -o ../certs/AmazonRootCA1.pem
for room in lounge kitchen bedroom1 bedroom2 bathroom; do ./add-device.sh $room; done
./add-device.sh nodered-gateway SmartLightingGatewayPolicy

echo "== 5. IoT rule -> both Lambdas"
aws iot create-topic-rule --rule-name LightingReadings --topic-rule-payload file://build/rule.json

echo "done :)"
```

infra/add-device.sh

```
#!/bin/bash
# makes one IoT "thing" + its own certificate, and attaches the policy
# usage: infra/add-device.sh <thingName> [policyName]
set -e
NAME=$1
POLICY=${2:-SmartLightingDevicePolicy}
CERTS="$(cd "$(dirname "$0")/.." && pwd)/certs"
```

```
aws iot create-thing --thing-name "$NAME" >/dev/null
ARN=$(aws iot create-keys-and-certificate --set-as-active \
  --certificate-pem-outfile "$CERTS/$NAME.cert.pem" \
  --private-key-outfile "$CERTS/$NAME.private.key" \
  --query certificateArn --output text)
aws iot attach-policy --policy-name "$POLICY" --target "$ARN"
aws iot attach-thing-principal --thing-name "$NAME" --principal "$ARN"
echo "  $NAME -> cert ${ARN##*/}"
```

infra/update-lambdas.sh

```
#!/bin/bash
# re-deploys the Lambda code after I change it
set -e
cd "$(dirname "$0")"; mkdir -p build
for pair in lighting-log-reading:log-reading lighting-control-light:control-light; do
  fn=${pair%:*}; dir=${pair##*:}
  (cd ../lambdas/$dir && zip -qj ../../infra/build/$dir.zip index.mjs)
  aws lambda update-function-code --function-name $fn --zip-file fileb://build/$dir.zip --query LastModified --
  output text
done
```

infra/teardown.sh

```
#!/bin/bash
# deletes everything setup.sh made (so I don't get billed after the unit's done)
cd "$(dirname "$0")"
export AWS_PAGER=""
aws iot delete-topic-rule --rule-name LightingReadings
# only my things (rooms, gateway, load-test rooms) - not anything else in the account
for t in $(aws iot list-things --query 'things[].thingName' --output text | tr '\t' '\n' | grep -E
'^(\lounge|kitchen|bedroom1|bedroom2|bathroom|nodered-gateway|load-[0-9]+)'); do
  for p in $(aws iot list-thing-principals --thing-name $t --query 'principals[]' --output text); do
    id=${p##*/}
    aws iot detach-thing-principal --thing-name $t --principal $p
    for pol in SmartLightingDevicePolicy SmartLightingGatewayPolicy; do aws iot detach-policy --policy-name $pol --
    target $p >/dev/null; done
    aws iot update-certificate --certificate-id $id --new-status INACTIVE
    aws iot delete-certificate --certificate-id $id
  done
  aws iot delete-thing --thing-name $t
done
aws iot delete-policy --policy-name SmartLightingDevicePolicy
aws iot delete-policy --policy-name SmartLightingGatewayPolicy
for fn in log control; do
  aws lambda delete-function --function-name lighting-${ [ $fn = log ] && echo log-reading || echo control-light}
  aws iam delete-role-policy --role-name lighting-$fn-role --policy-name lighting-$fn-access
  aws iam delete-role --role-name lighting-$fn-role
done
aws dynamodb delete-table --table-name SmartLighting >/dev/null
echo "all gone"
```

loadtest/provision-load-rooms.sh

```
#!/bin/bash
# makes load-001 ... load-N, each with its own cert (same as a real device would have)
# usage: loadtest/provision-load-rooms.sh 100
N=${1:-100}
cd "$(dirname "$0")/.."
seq -f "load-%03g" 1 $N | xargs -P 6 -I{} infra/add-device.sh {}
ls certs/load-*.cert.pem | wc -l | xargs echo "load room certs:"
```

infra/policies/device-policy.json

```
{
  "Version": "2012-10-17",
  "Statement": [
    {
      "Effect": "Allow",
      "Action": "iot:Connect",
      "Resource": "arn:aws:iot:REGION:ACCOUNT:client/${iot:Connection.Thing.ThingName}",
      "Condition": { "Bool": { "iot:Connection.Thing.IsAttached": "true" } }
    },
    {
      "Effect": "Allow",
      "Action": "iot:Publish",
```

```

    "Resource": "arn:aws:iot:REGION:ACCOUNT:topic/home/${iot:Connection.Thing.ThingName}/telemetry"
  },
  {
    "Effect": "Allow",
    "Action": "iot:Subscribe",
    "Resource": "arn:aws:iot:REGION:ACCOUNT:topicfilter/home/${iot:Connection.Thing.ThingName}/cmd"
  },
  {
    "Effect": "Allow",
    "Action": "iot:Receive",
    "Resource": "arn:aws:iot:REGION:ACCOUNT:topic/home/${iot:Connection.Thing.ThingName}/cmd"
  }
]
}

```

infra/policies/gateway-policy.json

```

{
  "Version": "2012-10-17",
  "Statement": [
    {
      "Effect": "Allow",
      "Action": "iot:Connect",
      "Resource": "arn:aws:iot:REGION:ACCOUNT:client/nodered-gateway",
      "Condition": { "Bool": { "iot:Connection.Thing.IsAttached": "true" } }
    },
    {
      "Effect": "Allow",
      "Action": "iot:Subscribe",
      "Resource": "arn:aws:iot:REGION:ACCOUNT:topicfilter/home/+/telemetry"
    },
    {
      "Effect": "Allow",
      "Action": "iot:Receive",
      "Resource": "arn:aws:iot:REGION:ACCOUNT:topic/home/*/telemetry"
    },
    {
      "Effect": "Allow",
      "Action": "iot:Publish",
      "Resource": [
        "arn:aws:iot:REGION:ACCOUNT:topic/home/*/clean",
        "arn:aws:iot:REGION:ACCOUNT:topic/home/alerts/*"
      ]
    }
  ]
}

```

infra/policies/log-role-policy.json

```

{
  "Version": "2012-10-17",
  "Statement": [
    {
      "Effect": "Allow",
      "Action": "dynamodb:UpdateItem",
      "Resource": "arn:aws:dynamodb:REGION:ACCOUNT:table/SmartLighting"
    },
    {
      "Effect": "Allow",
      "Action": ["logs:CreateLogGroup", "logs:CreateLogStream", "logs:PutLogEvents"],
      "Resource": "arn:aws:logs:REGION:ACCOUNT:log-group:/aws/lambda/lighting-log-reading:*"
    }
  ]
}

```

infra/policies/control-role-policy.json

```

{
  "Version": "2012-10-17",
  "Statement": [
    {
      "Effect": "Allow",
      "Action": ["dynamodb:Query", "dynamodb:GetItem", "dynamodb:UpdateItem"],
      "Resource": "arn:aws:dynamodb:REGION:ACCOUNT:table/SmartLighting"
    },
    {
      "Effect": "Allow",
      "Action": "iot:Publish",
      "Resource": "arn:aws:iot:REGION:ACCOUNT:topic/home/*/cmd"
    }
  ]
}

```

```

    "Effect": "Allow",
    "Action": ["logs:CreateLogGroup", "logs:CreateLogStream", "logs:PutLogEvents"],
    "Resource": "arn:aws:logs:REGION:ACCOUNT:log-group:/aws/lambda/lighting-control-light:*"
  }
}

```

infra/policies/lambda-trust.json

```

{
  "Version": "2012-10-17",
  "Statement": [
    { "Effect": "Allow", "Principal": { "Service": "lambda.amazonaws.com" }, "Action": "sts:AssumeRole" }
  ]
}

```

infra/policies/rule.json

```

{
  "sql": "SELECT * FROM 'home/+clean'",
  "awsIotSqlVersion": "2016-03-23",
  "ruleDisabled": false,
  "description": "cleaned readings from Node-RED -> log function + control function",
  "actions": [
    { "lambda": { "functionArn": "arn:aws:lambda:REGION:ACCOUNT:function:lighting-log-reading" } },
    { "lambda": { "functionArn": "arn:aws:lambda:REGION:ACCOUNT:function:lighting-control-light" } }
  ]
}

```

package.json

```

{
  "name": "sit314_distinction",
  "version": "1.0.0",
  "main": "index.js",
  "directories": {
    "doc": "docs"
  },
  "scripts": {
    "test": "echo \"Error: no test specified\" && exit 1"
  },
  "keywords": [],
  "author": "",
  "license": "ISC",
  "description": "",
  "dependencies": {
    "@aws-sdk/client-dynamodb": "^3.1146.0",
    "@aws-sdk/lib-dynamodb": "^3.1146.0",
    "dotenv": "^18.0.5",
    "express": "^5.2.1",
    "mqtt": "^5.16.0"
  }
}

```

Appendix D – Raw Output

evidence/loadtest_console.log (load test terminal output, trimmed – the full file has every light switch)

```
load test: steps 25 -> 50 -> 100 rooms, 120s each, every room sends every 1000ms, 4 worker processes
started 2026-10-04T10:59:55.078Z
t=10s step 1/3 (25 rooms)
... (first few of the thousands of light switch lines) ...
[21:59:56] load-002 >>> light switched ON (dark (avg 35 < 250 lux) + motion) latency=316ms
[21:59:57] load-004 >>> light switched ON (dark (avg 0 < 250 lux) + motion) latency=341ms
[21:59:57] load-005 >>> light switched ON (dark (avg 1 < 250 lux) + motion) latency=242ms
[21:59:57] load-001 >>> light switched ON (dark (avg 24 < 250 lux) + motion) latency=313ms
[21:59:57] load-003 >>> light switched ON (dark (avg 21 < 250 lux) + motion) latency=297ms
[21:59:58] load-009 >>> light switched ON (dark (avg 0 < 250 lux) + motion) latency=345ms
...
t=20s step 1/3 (25 rooms)
t=30s step 1/3 (25 rooms)
t=40s step 1/3 (25 rooms)
t=50s step 1/3 (25 rooms)
t=60s step 1/3 (25 rooms)
t=70s step 1/3 (25 rooms)
t=80s step 1/3 (25 rooms)
t=90s step 1/3 (25 rooms)
t=100s step 1/3 (25 rooms)
t=110s step 1/3 (25 rooms)
t=120s step 2/3 (50 rooms)
t=130s step 2/3 (50 rooms)
t=140s step 2/3 (50 rooms)
t=150s step 2/3 (50 rooms)
t=160s step 2/3 (50 rooms)
t=170s step 2/3 (50 rooms)
t=180s step 2/3 (50 rooms)
t=190s step 2/3 (50 rooms)
t=200s step 2/3 (50 rooms)
t=210s step 2/3 (50 rooms)
t=220s step 2/3 (50 rooms)
t=230s step 2/3 (50 rooms)
t=240s step 3/3 (100 rooms)
t=250s step 3/3 (100 rooms)
t=260s step 3/3 (100 rooms)
t=270s step 3/3 (100 rooms)
t=280s step 3/3 (100 rooms)
t=290s step 3/3 (100 rooms)
t=300s step 3/3 (100 rooms)
t=310s step 3/3 (100 rooms)
t=320s step 3/3 (100 rooms)
t=330s step 3/3 (100 rooms)
t=340s step 3/3 (100 rooms)
t=350s step 3/3 (100 rooms)
t=360s step 3/3 (100 rooms)
```

=== results ===

(index)	rooms	msgs sent	msg/s	cmds back	p50 ms	p95 ms	p99 ms	max ms
0	25	2925	24.4	46	314	813	1449	1449
1	50	5920	49.3	353	1462	43511	60744	65824
2	100	11842	98.7	1333	22411	83839	103100	120284

total sent 20687, commands received 1732, mqtt errors 0
saved /Users/adityasuhag/sit314_distinction/loadtest/./evidence/loadtest_2026-10-04T1059.json

evidence/security_test.log

```
[21:49:25] --- control: lounge cert, clientId lounge, publish to home/lounge/telemetry
[21:49:26] connected
[21:49:26] result: publish accepted (PUBACK received) – expected
[21:49:26] --- test 1: lounge cert publishes to home/kitchen/telemetry
[21:49:26] connected
[21:49:26] result: connection closed by AWS
[21:49:27] --- test 2: lounge cert connects as clientId "kitchen"
[21:49:27] result: connection closed by AWS
[21:49:27] --- test 3: no client certificate
[21:49:27] result: error -> read ECONNRESET
```

evidence/setup.log (running infra/setup.sh on a fresh account)

```
account 789633759313, region ap-southeast-2, endpoint a25bw3rao5ztzni-ats.iot.ap-southeast-2.amazonaws.com
== 1. DynamoDB table
== 2. IAM roles (one each, only what that function needs)
    waiting for IAM to catch up...
== 3. Lambda functions
    lighting-log-reading deployed
    lighting-control-light deployed
== 4. IoT Core policies, things and certificates
    lounge -> cert ba6348a9c1a08f929b6b62a05c6f592372a353f10ff1ed1930b5fba30d33f54b
    kitchen -> cert 7d73e906d87f93b6160d829e10a715224dd03b583d787fd311c737e626a3fb69
    bedroom1 -> cert 1d0891202a04be0ec4598e4264ae84338304d94dea13a70877d571b74c2b3d60
    bedroom2 -> cert 1e2d7b38b3257c0599085d3ceedafc48d55064ca12453951995d9943ac58f031
    bathroom -> cert 7b4c7e4e86b2d2ee5deab38593529968903674a716e79febfdeea29bbb42a825
    nodered-gateway -> cert c1717d9f0eaf2d9084695ee2a1cdccfcdf47aebd8b8720b50240f9c9cec3a5a
== 5. IoT rule -> both Lambdas
done :)
```